

# Timing and Synchronization

## Chapter 14

- During a baseball game, I receive two messages:
- “The hitter hits a homerun!”
- “Timeout! Pitcher A is replaced by Pitcher B”
- Q: Who surrenders the homerun? A or B?

- During a baseball game, I receive two messages:
- “The hitter hits a homerun!”
  - 19:00:03
- “Timeout! Pitcher A is replaced by Pitcher B”
  - 19:00:02
- Q: Who surrenders the homerun? A or B?
- Sol 1: Check the time stamps of the messages
- Only works when clocks are synchronized

# Coordinated Universal Time (UTC)



- Astronomical time: Historically, time is defined by the relative positions between the earth and the sun
- UTC: An authoritative atomic clock with very high precision
  - “Leap second” is added occasionally to make it consistent with astronomical time
  - Used to synchronize all satellites
- GPS satellites broadcast time information to land-based devices
  - Precision is within 0.1 - 10ms from UTC
- GPS time is not always available

# Some Definitions

- An Asynchronous Distributed System consists of a number of **processes**.
- Each process has a **state** (values of variables).
- Each process takes **actions** to change its state, which may be an **instruction** or a communication action (**send**, **receive**).
- An **event** is the occurrence of an action.
- Each process has a local clock – events *within* a process can be assigned **timestamps**, and thus ordered linearly.
- But – in a distributed system, we also need to know the time order of events across different processes.

# Clocks

- Hardware clock: each computer has a device counting the oscillations of a crystal
- Denote  $H_i(t)$  as the hardware clock of process  $i$
- The operating system translates the hardware clock into a software clock:
- $C_i(t) = \alpha H_i(t) + \beta$
- Ideally, we want  $C_i(t) = t$
- To do clock synchronization, OS changes the values of  $\alpha$  and  $\beta$

# Clock Skew and Clock Drift

- Each process (running at some end host) has its own clock.
- When comparing two clocks at two processes:
  - Clock **Skew** = Relative Difference in clock *values* of two processes
    - Like distance between two vehicles on a road
    - The error in  $\beta$
  - Clock **Drift** = Relative Difference in clock *frequencies (rates)* of two processes
    - Like difference in speeds of two vehicles on the road
    - The error in  $\alpha$
- A non-zero clock skew implies clocks are not synchronized.
- A non-zero clock drift causes skew to increase (eventually).
- The skew of a typically computer is about  $10^{-6}$ sec/sec
- About 3ms error per hour

# Types of Synchronization

- Consider a group of processes
- External Synchronization
  - Each process  $i$ 's clock is within a bound  $D$  of a well-known clock  $S(t)$  external to the group
  - $|C_i(t) - S(t)| < D$  at all times
  - External clock may be connected to UTC (Universal Coordinated Time) or an atomic clock
- Internal Synchronization
  - Every pair of processes in group have clocks within bound  $D$
  - $|C_i(t) - C_j(t)| < D$  at all times and for all processes  $i, j$
- External Synchronization with  $D \Rightarrow$  Internal Synchronization with  $2 \cdot D$ 
  - Why?
- Internal Synchronization does not imply External Synchronization



# Correctness of Hardware Clocks

- Require that the drift of a hardware clock cannot exceed some threshold,  $\rho$
- In other words, given  $t' > t$
- $(1 - \rho)(t' - t) \leq H(t') - H(t) \leq (1 + \rho)(t' - t)$

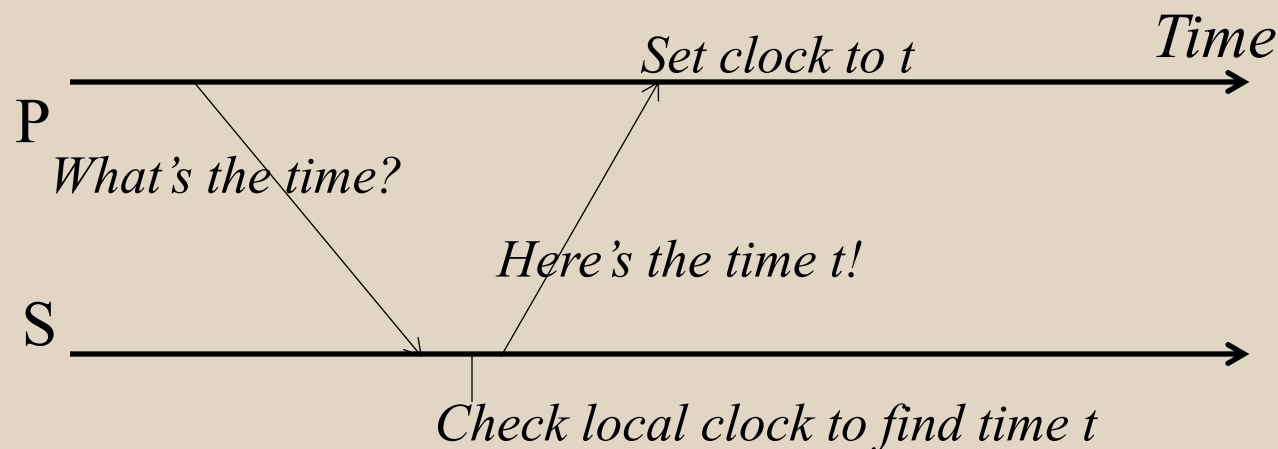
# Monotonicity of Software Clocks

- For causality, it is sometimes required that  $C_i(t)$  needs to be non-decreasing
  - The timestamp of “result” is always larger than the timestamp of “cause”
- During synchronization, process  $i$  obtains UTC time  $S(t)$
- If  $S(t) > C_i(t)$ , set  $C_i(t) = S(t)$
- If  $S(t) < C_i(t)$ , we cannot set  $C_i(t) = S(t)$
- Recall  $C_i(t) = \alpha H_i(t) + \beta$
- Reduce  $\alpha$  for some time

# CRISTIAN'S ALGORITHM

# Basics

- External time synchronization
- All processes  $P$  synchronize with a time server  $S$

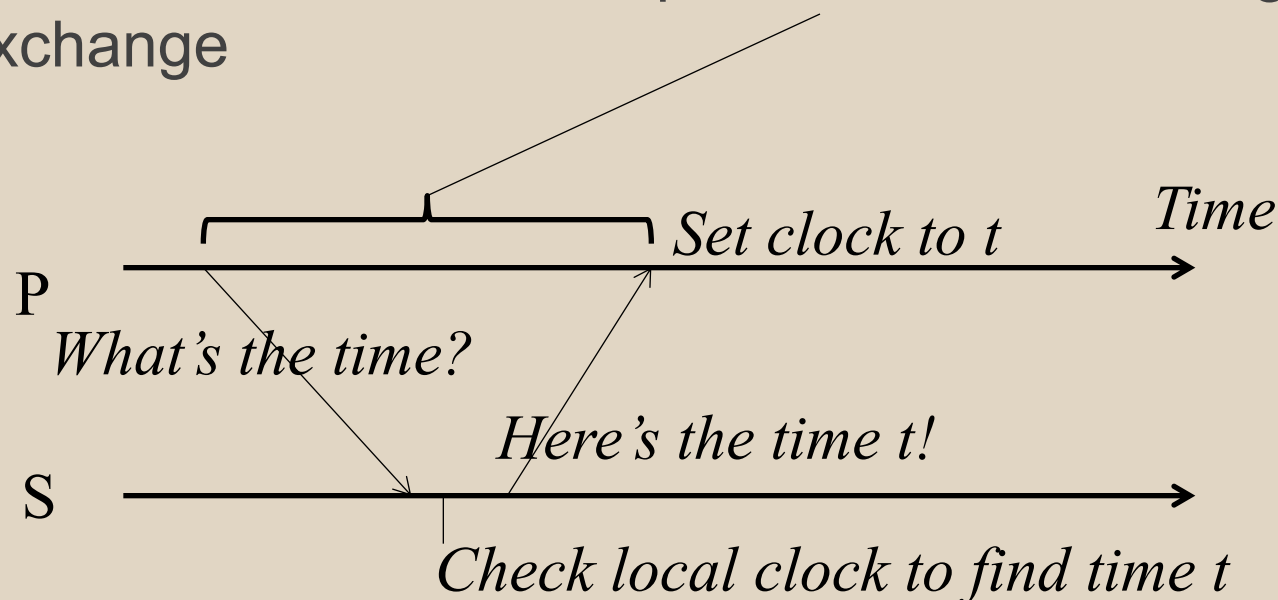


# What's Wrong

- By the time response message is received at P, time has moved on
- P's time set to  $t$  is inaccurate!
- Inaccuracy a function of message latencies
- Since latencies unbounded in an asynchronous system, the inaccuracy cannot be bounded


# Cristian's Algorithm

- P measures the round-trip-time RTT of message exchange



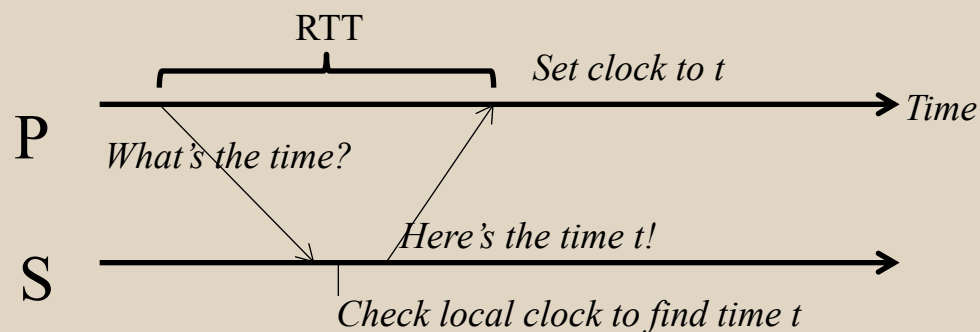
# Cristian's Algorithm (2)

- P measures the round-trip-time RTT of message exchange
- Suppose we know the minimum  $P \rightarrow S$  latency  $\text{min1}$
- And the minimum  $S \rightarrow P$  latency  $\text{min2}$ 
  - $\text{min1}$  and  $\text{min2}$  depend on Operating system overhead to buffer messages, TCP time to queue messages, etc.



# Cristian's Algorithm (3)


- P measures the round-trip-time RTT of message exchange
- Suppose we know the minimum  $P \rightarrow S$  latency  $\min1$
- And the minimum  $S \rightarrow P$  latency  $\min2$ 
  - $\min1$  and  $\min2$  depend on Operating system overhead to buffer messages, TCP time to queue messages, etc.
- The actual time at P when it receives response is between  $[t + \min2, t + \text{RTT} - \min1]$





# Cristian's Algorithm (4)


- The actual time at P when it receives response is between  $[t + \min_2, t + \text{RTT} - \min_1]$
- P sets its time to halfway through this interval
  - To:  $t + (\text{RTT} + \min_2 - \min_1) / 2$
- Error is at most  $(\text{RTT} - \min_2 - \min_1) / 2$ 
  - Bounded!



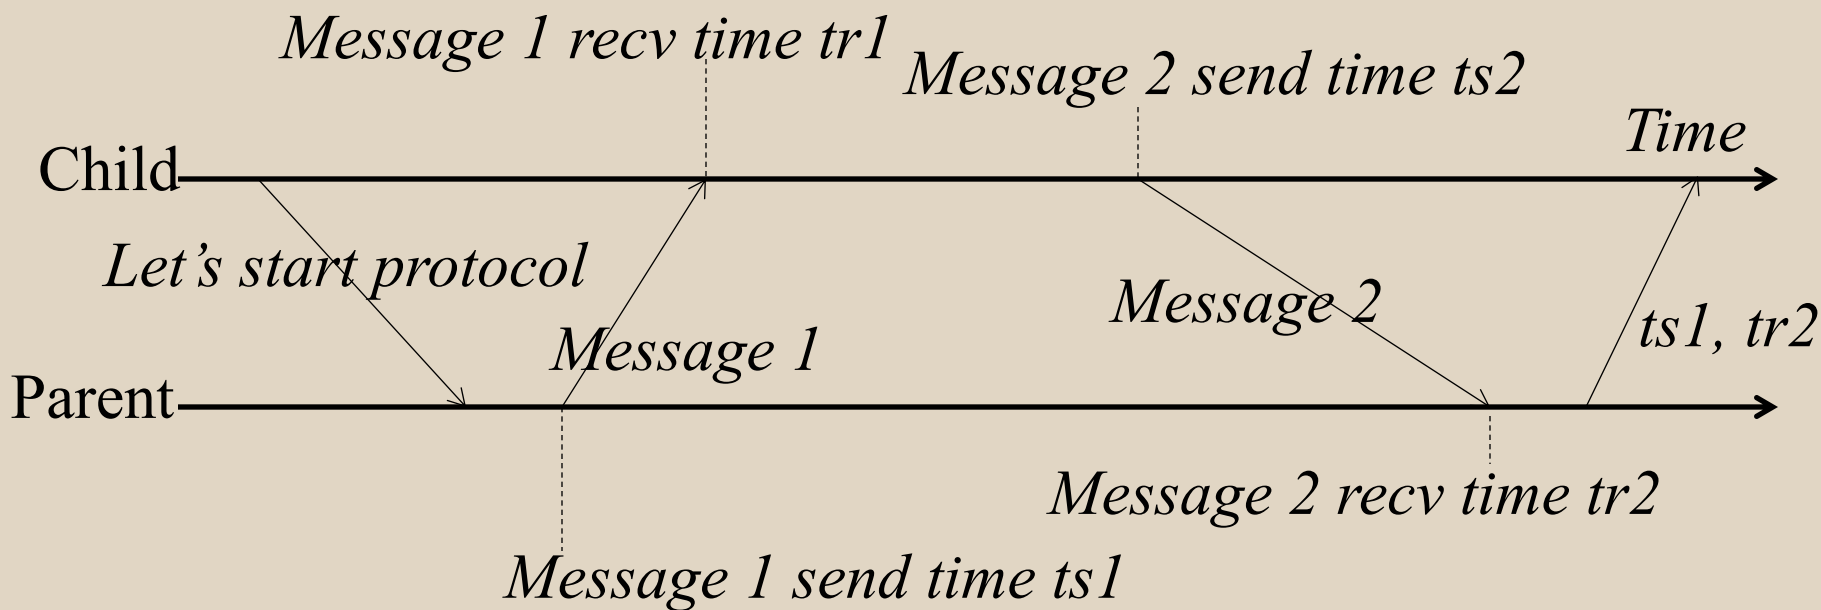
# NETWORK TIME PROTOCOL (NTP)

# NTP = Network Time Protocol

- NTP Servers organized in a tree
- Each Client = a leaf of tree
- Each node synchronizes with its tree parent



# NTP Protocol



# Basic Properties

- Let  $t$  and  $t'$  be the transmission times for Message 1 and Message 2, respectively
- Let  $d$  be the total transmission time of the two messages
  - $d = t + t'$
- Let  $o$  be the offset of the clocks
- We now have
- $tr1 = ts1 + t + o$
- $tr2 = ts2 + t' - o$
- $d = t + t' = tr1 - ts1 + tr2 - ts2$  (Can be thought of as RTT)
- $o = o_i + (t' - t)/2$ , where  $o_i = (tr1 - ts1 + ts2 - tr2)/2$

# What the Child Does

- Child calculates *offset* between its clock and parent's clock
- Uses  $ts1, tr1, ts2, tr2$
- Offset is estimated as
 
$$o \approx o_i = (tr1 - tr2 + ts2 - ts1)/2$$
- Error is bounded by  $d/2$
- NTP uses multiple pairs of  $(o_i, d)$  to obtain a more accurate clock

# And yet...

- **We still have a non-zero error!**
- **We just can't seem to get rid of error**
  - Can't, as long as message latencies are non-zero
- **Can we avoid synchronizing clocks altogether, and still be able to order events?**

# LOGICAL CLOCK



# Ordering Events in a Distributed System

- To order events across processes, trying to sync clocks is one approach
- What if we instead assigned timestamps to events that were not *absolute* time?
- As long as these timestamps obey *causality*, that would work

If an event A causally happens before another event B, then  $\text{timestamp}(A) < \text{timestamp}(B)$

Humans use causality all the time

E.g., I enter a house only after I unlock it

E.g., You receive a letter only after I send it


# Logical (or Lamport) Ordering

- Proposed by Leslie Lamport in the 1970s
- Used in almost all distributed systems since then
- Almost all cloud computing systems use some form of logical ordering of events


# Logical (or Lamport) Ordering(2)

- Define a logical relation *Happens-Before* among pairs of events
- *Happens-Before* denoted as  $\rightarrow$
- Three rules
  1. On the same process:  $a \rightarrow b$ , if  $time(a) < time(b)$  (using the local clock)
  2. If p1 sends  $m$  to p2:  $send(m) \rightarrow receive(m)$
  3. (Transitivity) If  $a \rightarrow b$  and  $b \rightarrow c$  then  $a \rightarrow c$
- Creates a *partial order* among events
  - Not all events related to each other via  $\rightarrow$

# Example



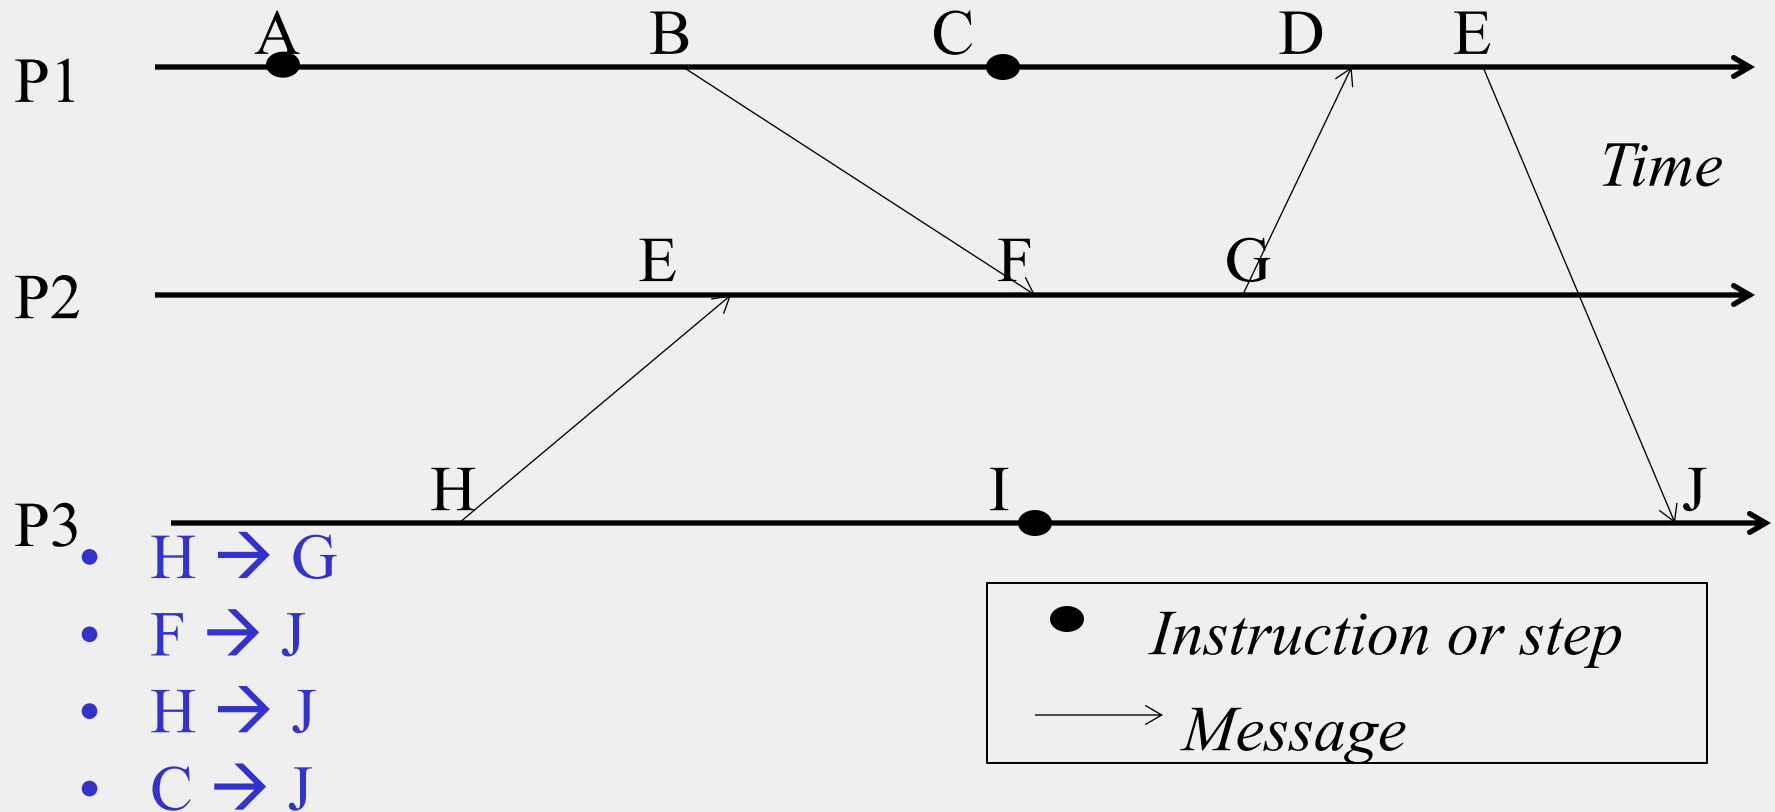
# Happens-Before



- $A \rightarrow B$
- $B \rightarrow F$
- $A \rightarrow F$




# Happens-Before (2)



# In practice: Lamport timestamps

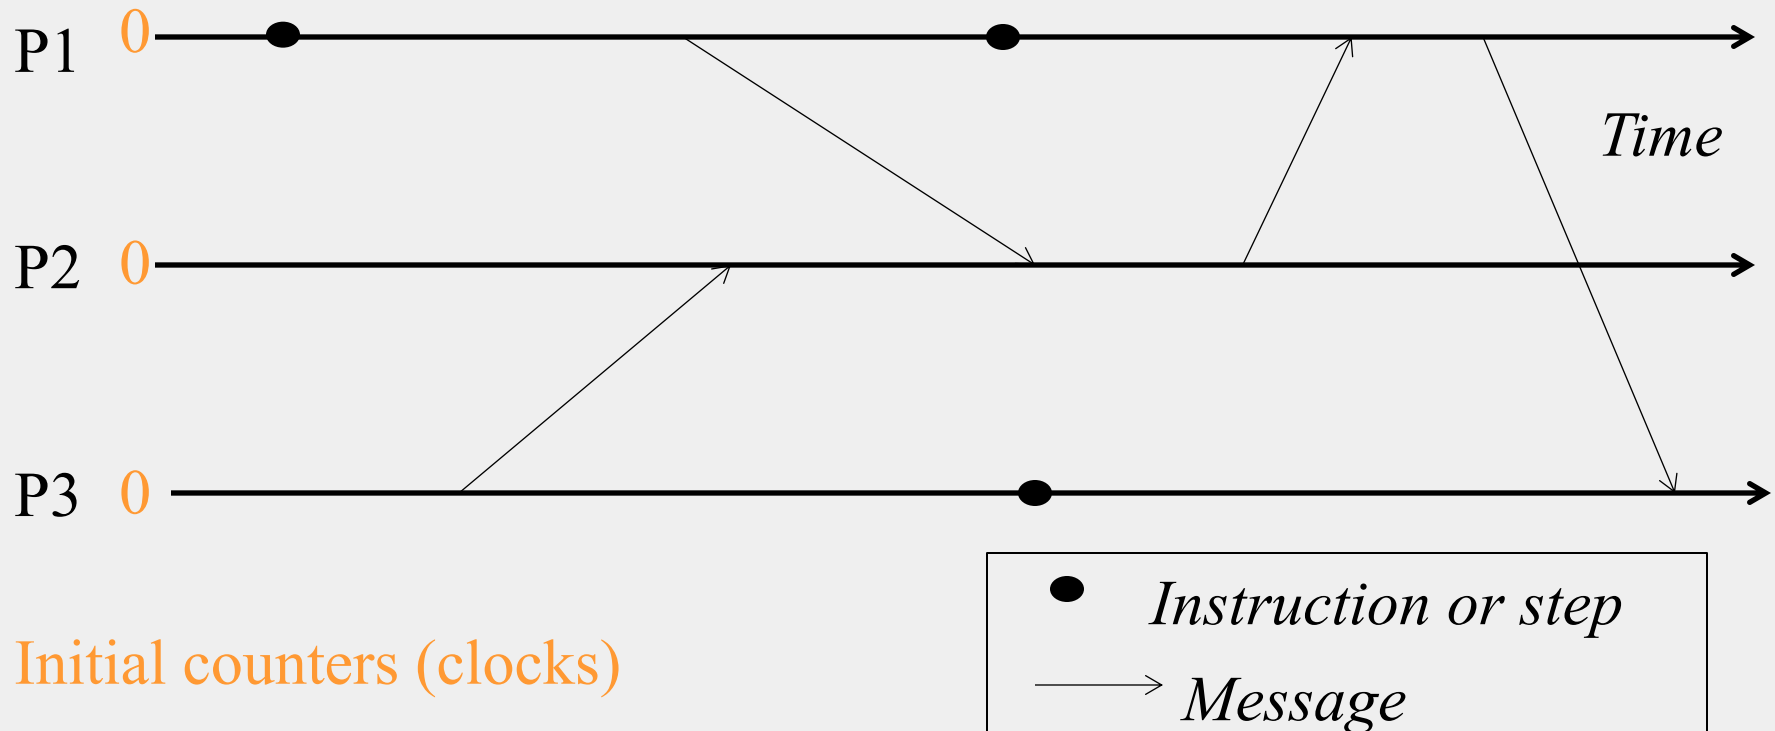
- **Goal: Assign logical (Lamport) timestamp to each event**
- **Timestamps obey causality**
- **Rules**
  - Each process uses a local counter (clock) which is an integer
    - initial value of counter is zero
  - A process increments its counter when a **send** or an **instruction** happens at it. The counter is assigned to the event as its timestamp.
  - A **send (message)** event carries its timestamp
  - For a **receive (message)** event the counter is updated by
$$\max(\text{local clock}, \text{message timestamp}) + 1$$

# Example

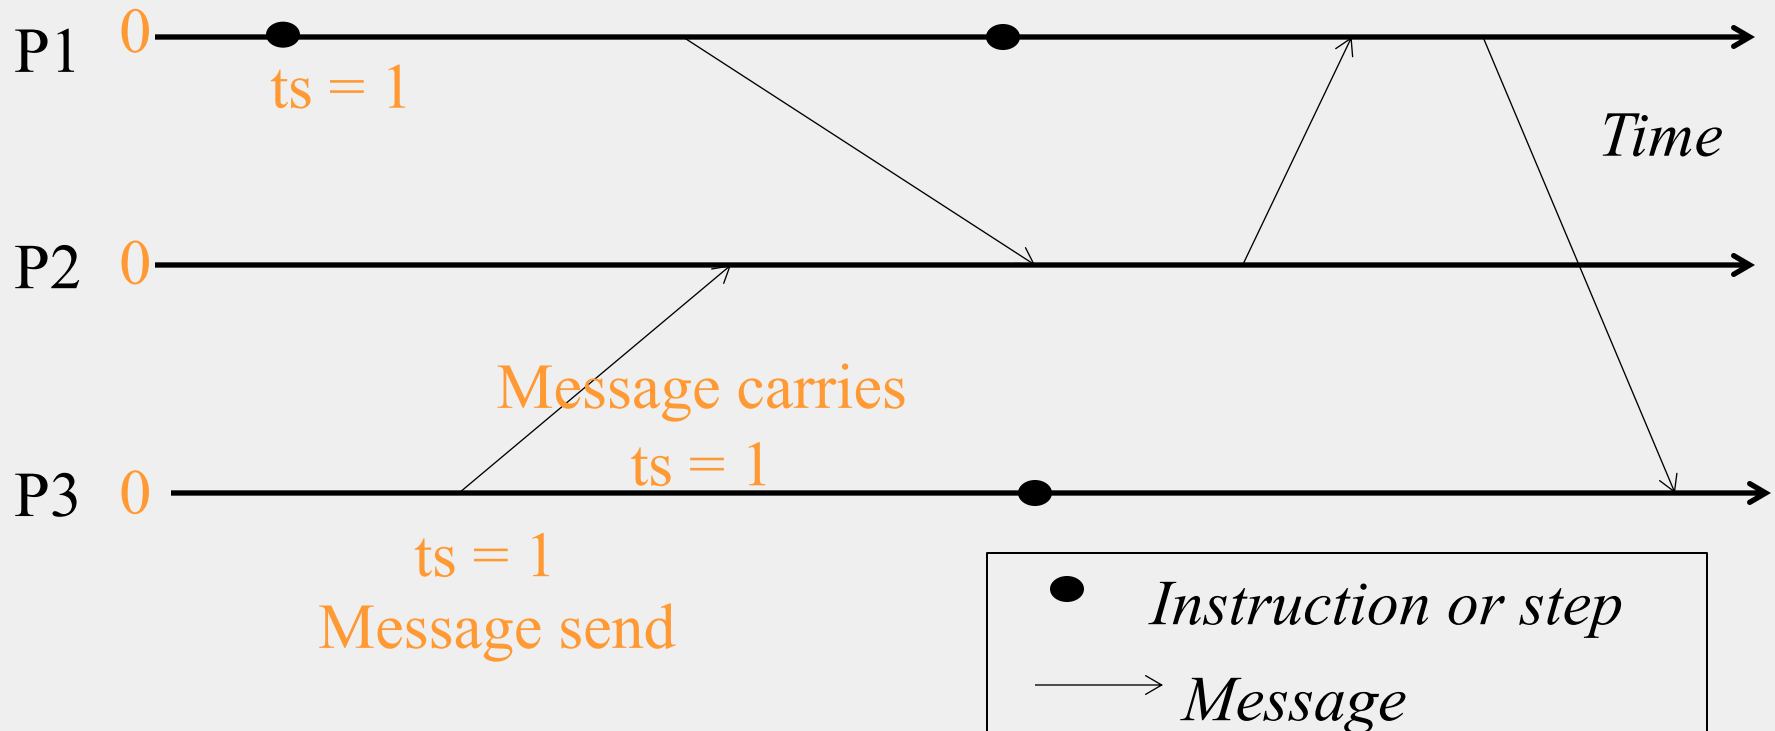




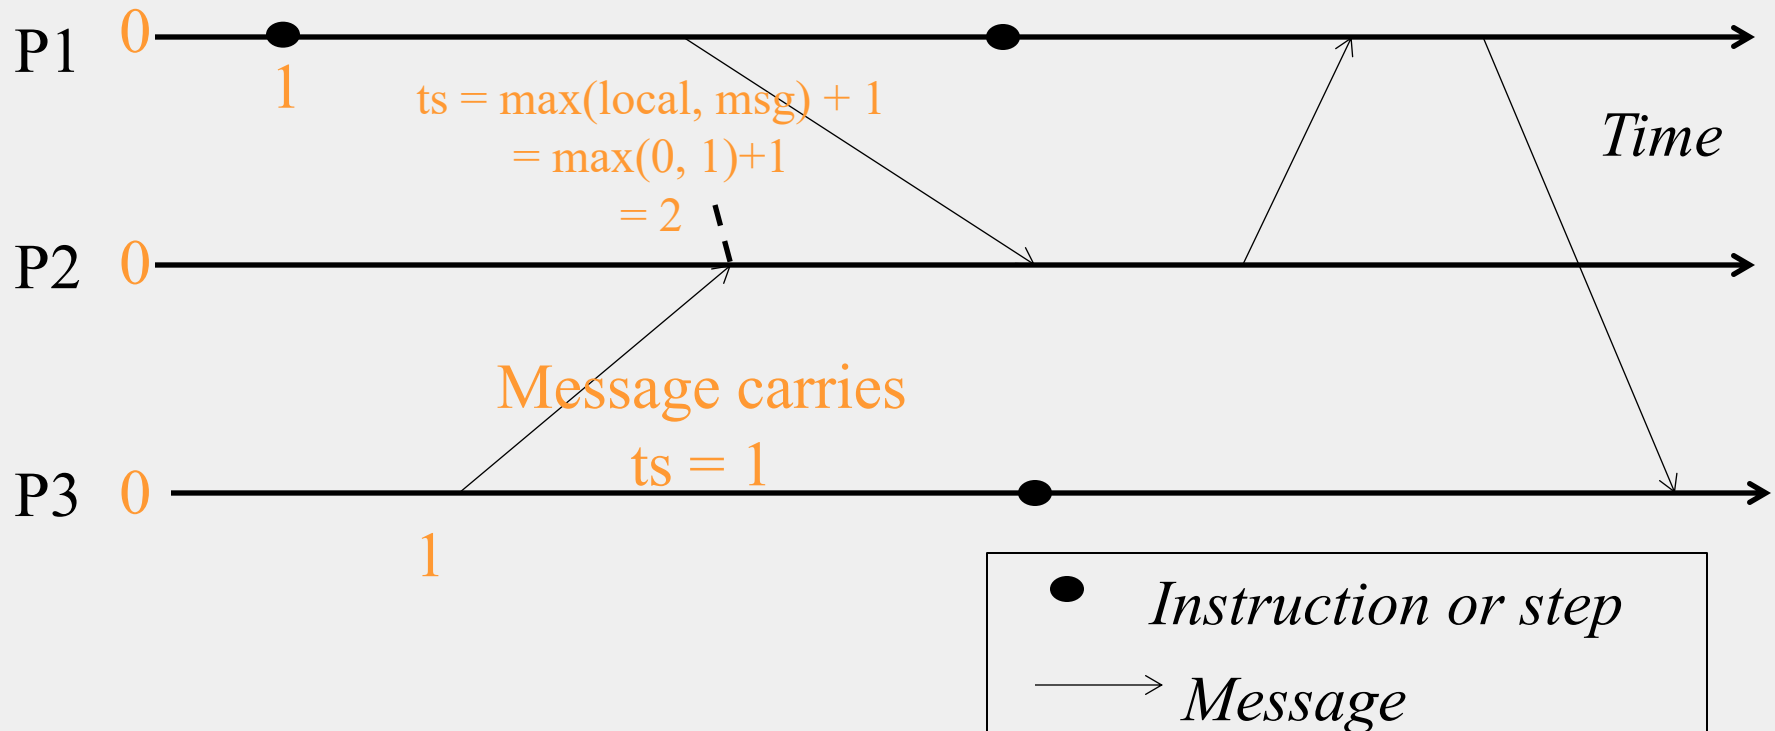
# Lamport Timestamps



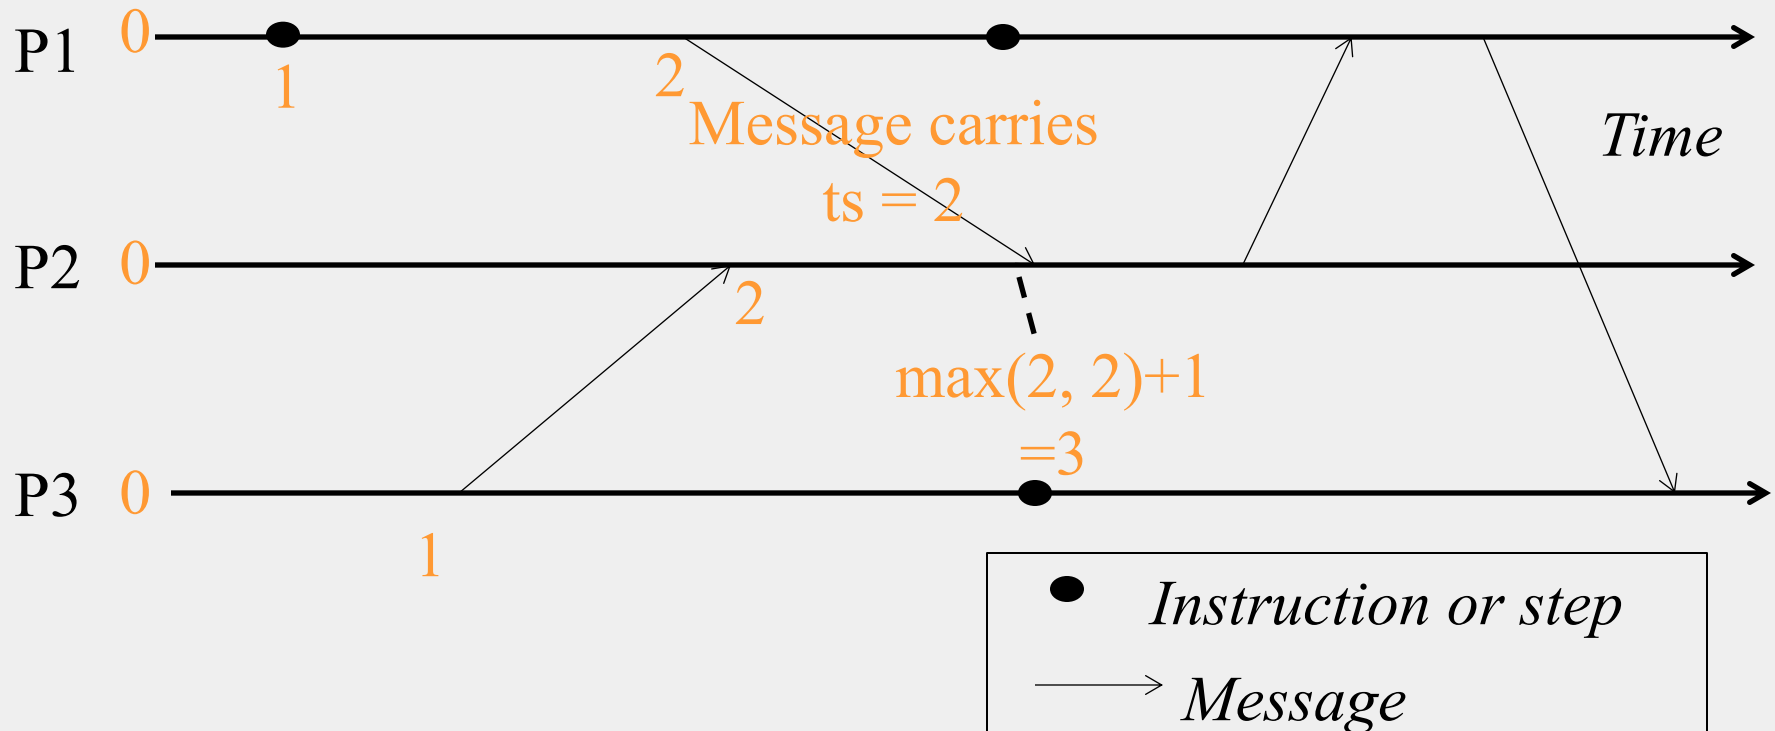
# Lamport Timestamps



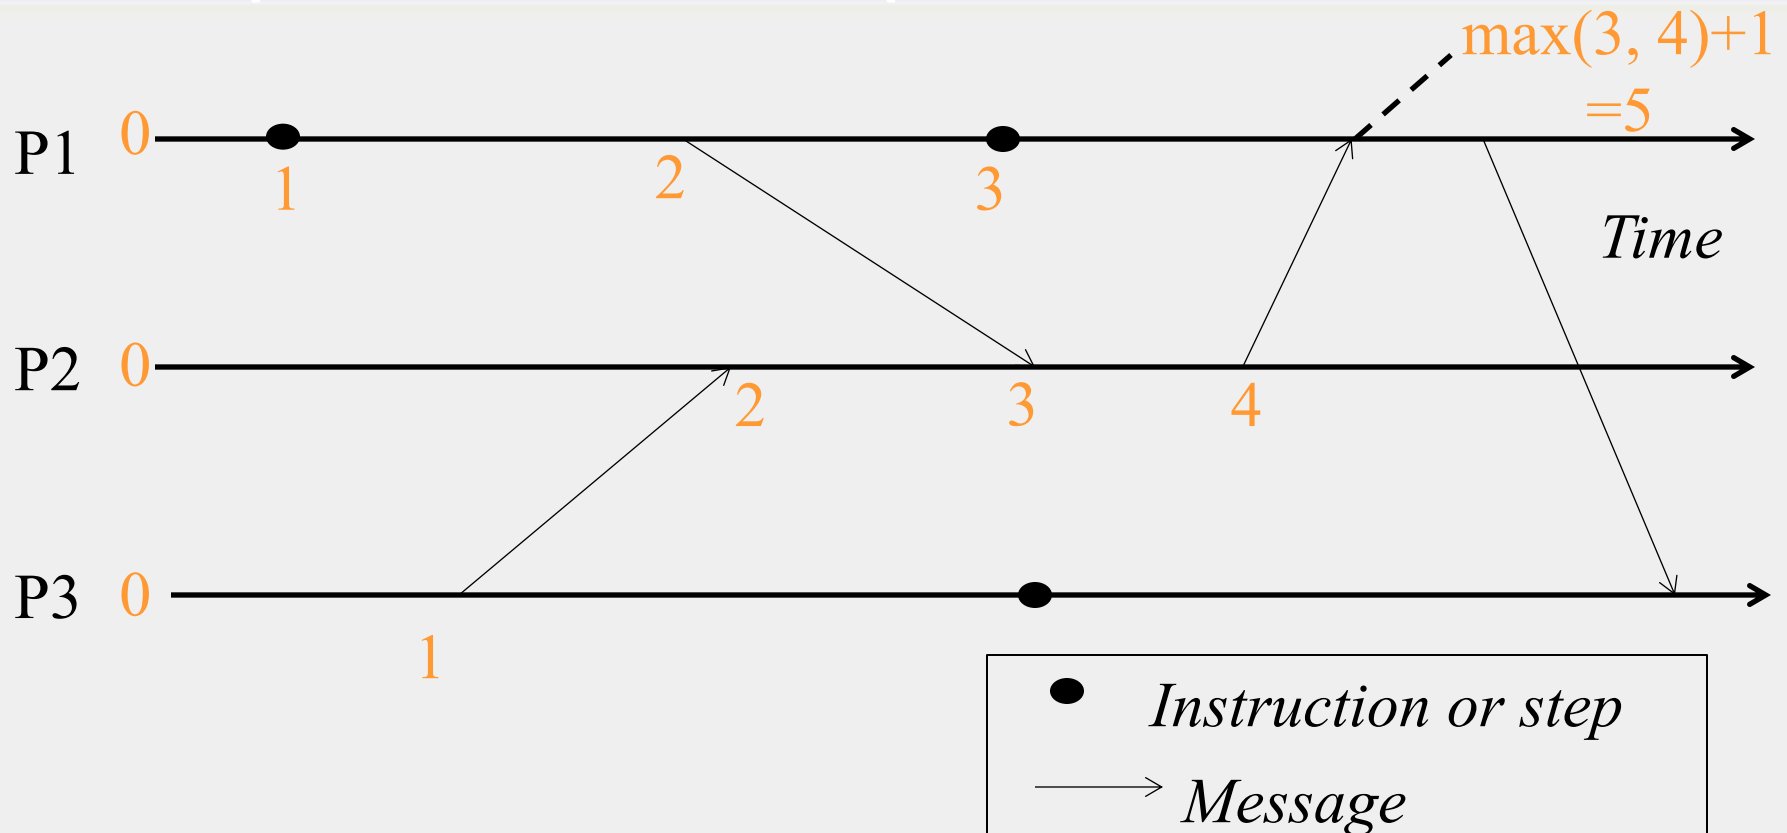
# Lamport Timestamps



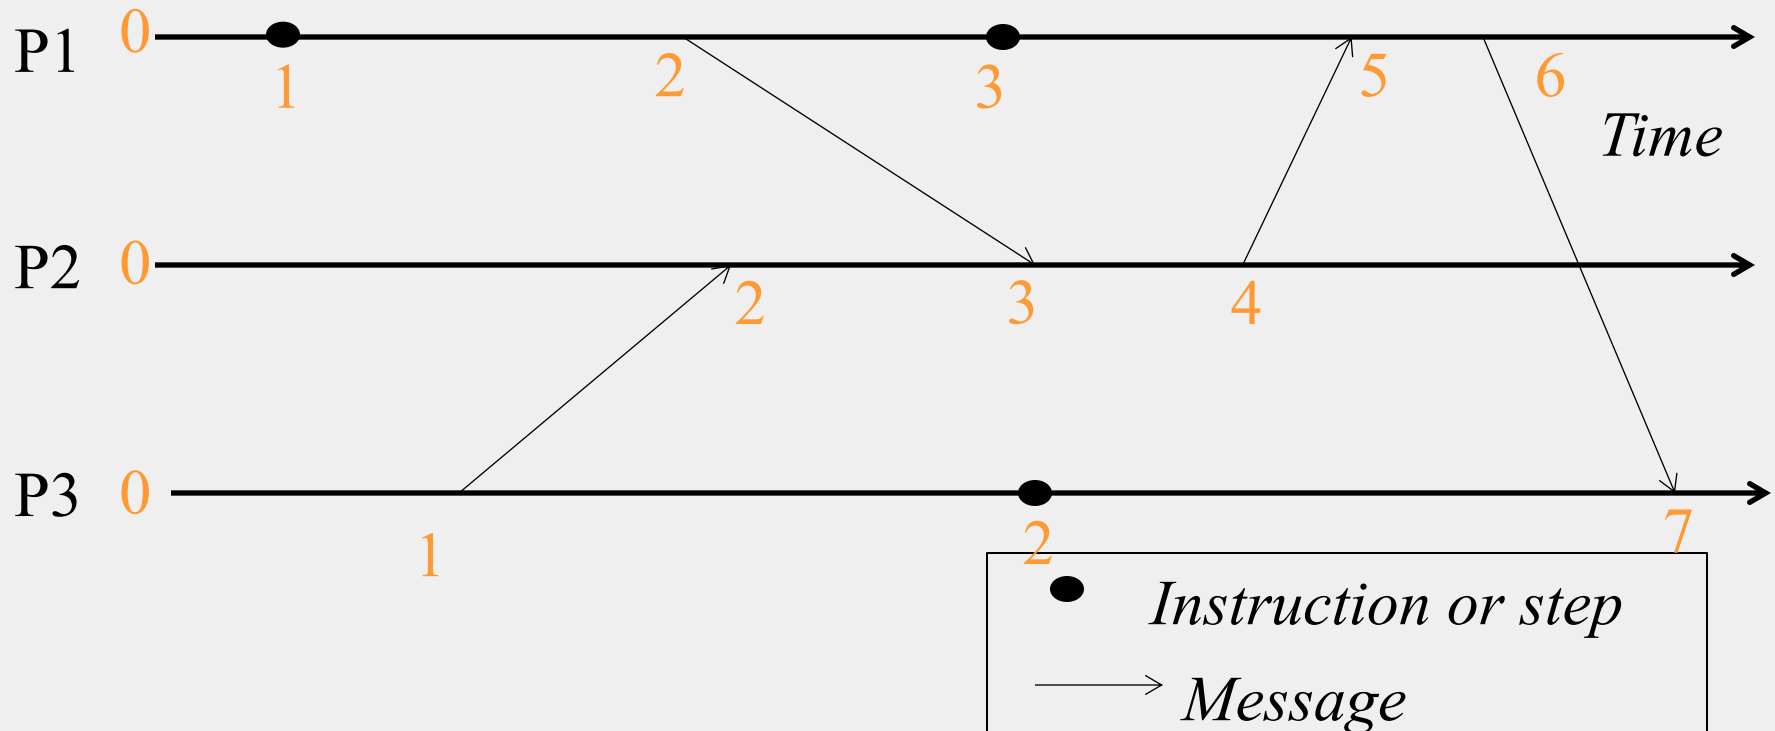
# Lamport Timestamps




# Lamport Timestamps



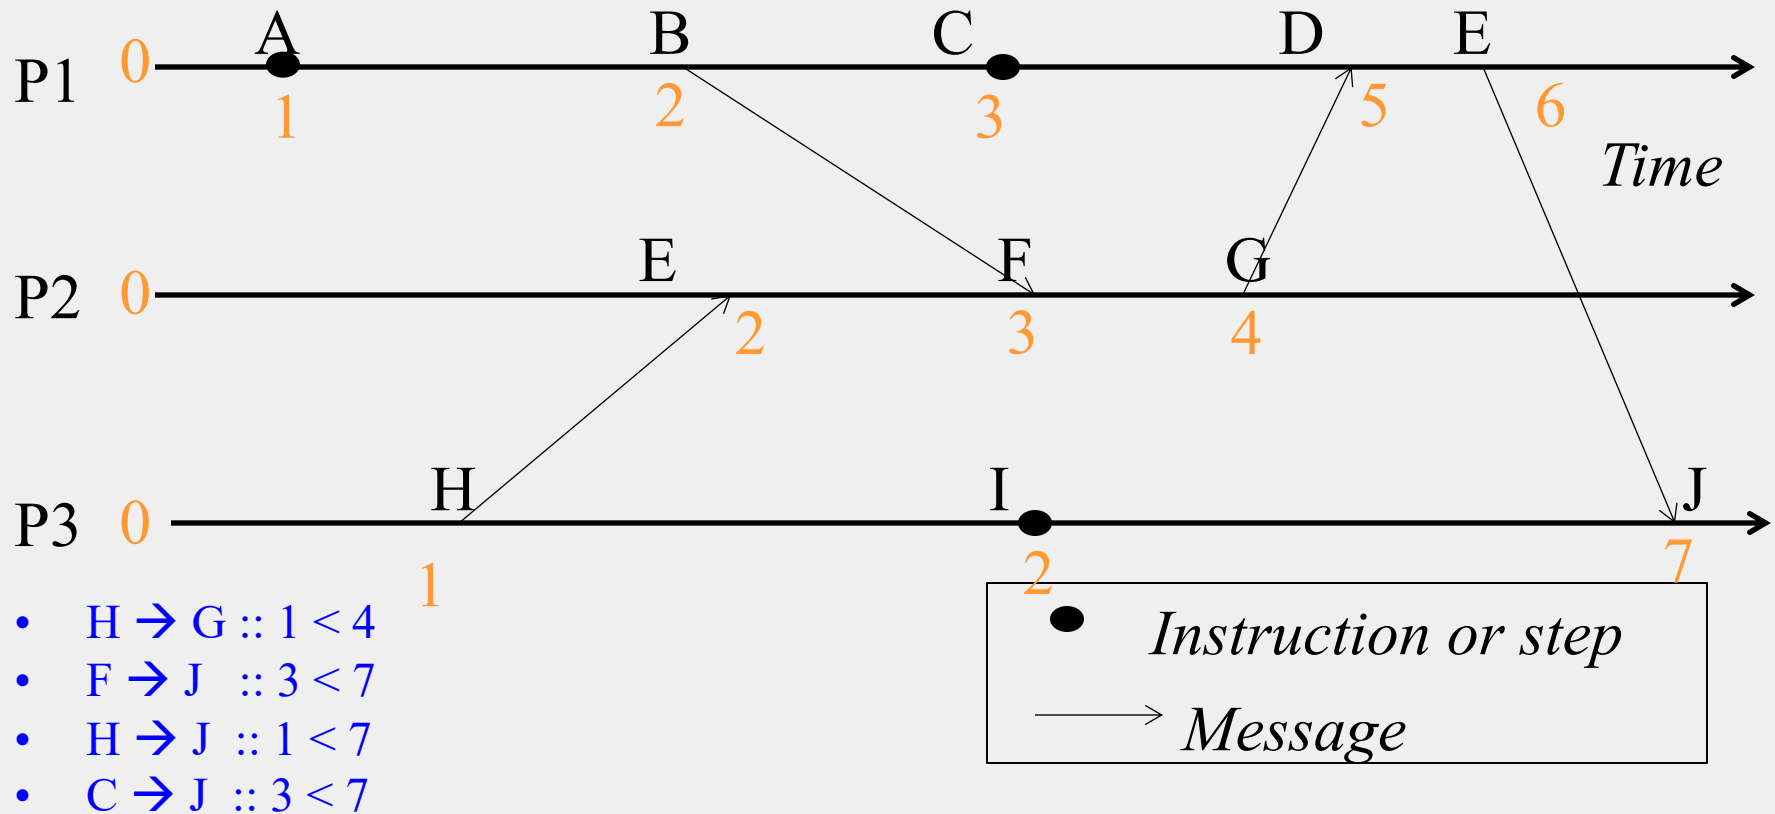
# Lamport Timestamps



# Obeying Causality




# Obeying Causality (2)





# Not always *implying* Causality



- $? C \rightarrow F ? :: 3 = 3$
- $? H \rightarrow C ? :: 1 < 3$
- (C, F) and (H, C) are pairs of concurrent events

# Concurrent Events

- **A pair of concurrent events doesn't have a causal path from one event to another (either way, in the pair)**
- **Lamport timestamps not guaranteed to be ordered or unequal for concurrent events**
- **Ok, since concurrent events are not causality related!**
- **Remember**

$E1 \rightarrow E2 \Rightarrow \text{timestamp}(E1) < \text{timestamp}(E2)$ , **BUT**

$\text{timestamp}(E1) < \text{timestamp}(E2) \Rightarrow$

$\{E1 \rightarrow E2\} \text{ OR } \{E1 \text{ and } E2 \text{ concurrent}\}$

# VECTOR CLOCK

# Vector Timestamps

- Used in key-value stores like Riak
- Each process uses a vector of integer clocks
- Suppose there are  $N$  processes in the group  $1 \dots N$
- Each vector has  $N$  elements
- Process  $i$  maintains vector  $V_i[1 \dots N]$
- $j$ th element of vector clock at process  $i$ ,  $V_i[j]$ , is  $i$ 's knowledge of latest events at process  $j$


# Assigning Vector Timestamps

- Incrementing vector clocks
  1. On an instruction or send event at process  $i$ , it increments only its  $i$ th element of its vector clock
  2. Each message carries the send-event's vector timestamp  $V_{\text{message}}[1 \dots N]$
  3. On receiving a message at process  $i$ :


$$V_i[i] = V_i[i] + 1$$

$$V_i[j] = \max(V_{\text{message}}[j], V_i[j]) \text{ for } j \neq i$$

# Example




# Vector Timestamps




Initial counters (clocks)

# Vector Timestamps

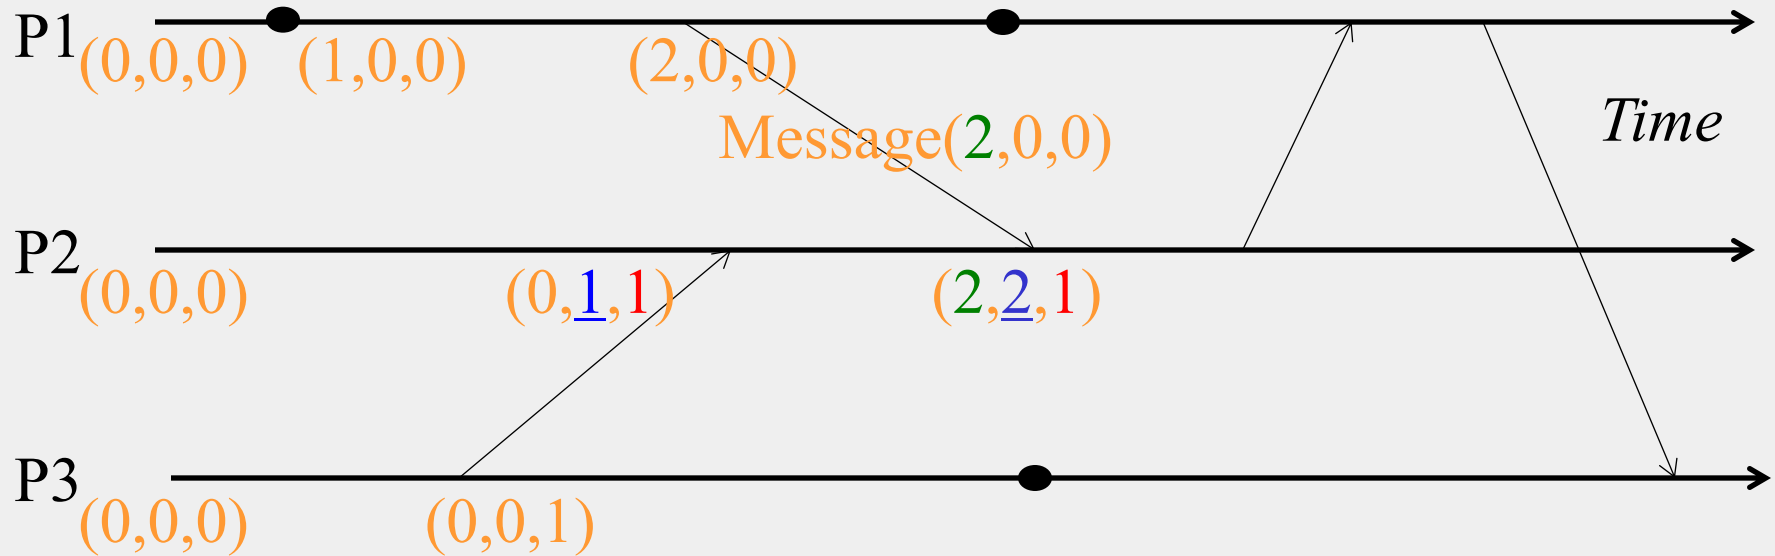





# Vector Timestamps



# Vector Timestamps



# Vector Timestamps



# Causally-Related ...

- $VT_1 = VT_2$ ,  
*iff* (if and only if)  
 $VT_1[i] = VT_2[i]$ , for all  $i = 1, \dots, N$
- $VT_1 \leq VT_2$ ,  
*iff*  $VT_1[i] \leq VT_2[i]$ , for all  $i = 1, \dots, N$
- Two events are **causally related** *iff*  
 $VT_1 < VT_2$ , i.e.,  
*iff*  $VT_1 \leq VT_2$  &  
there exists  $j$  such that  
 $1 \leq j \leq N$  &  $VT_1[j] < VT_2[j]$

## ... or Not Causally-Related


- Two events  $VT_1$  and  $VT_2$  are **concurrent**

*iff*

$$\text{NOT } (VT_1 \leq VT_2) \text{ AND NOT } (VT_2 \leq VT_1)$$


We'll denote this as  $VT_2 \parallel VT_1$

# Obeying Causality




- $A \rightarrow B :: (1,0,0) < (2,0,0)$
- $B \rightarrow F :: (2,0,0) < (2,2,1)$
- $A \rightarrow F :: (1,0,0) < (2,2,1)$

# Obeying Causality (2)



- $H \rightarrow G :: (0, 0, 1) < (2, 3, 1)$
- $F \rightarrow J :: (2, 2, 1) < (5, 3, 3)$
- $H \rightarrow J :: (0, 0, 1) < (5, 3, 3)$
- $C \rightarrow J :: (3, 0, 0) < (5, 3, 3)$

# Identifying Concurrent Events



- C & F :: (3,0,0) ||| (2,2,1)
- H & C :: (0,0,1) ||| (3,0,0)
- (C, F) and (H, C) are pairs of concurrent events



# Logical Timestamps: Summary

- **Lamport timestamps**
  - Integer clocks assigned to events
  - Obey causality
  - Cannot distinguish concurrent events
- **Vector timestamps**
  - Obey causality
  - By using more space, can also identify concurrent events

# Time and Ordering: Summary

- **Clocks are unsynchronized in an asynchronous distributed system**
- **But need to order events, across processes!**
- **Time synchronization**
  - Cristian's algorithm
  - NTP
  - Berkeley algorithm
  - But error a function of round-trip-time
- **Can avoid time sync altogether by instead assigning logical timestamps to events**